Evidence-Based Philosophy: A Foundational Methodological Framework for the Transformation of Philosophical Inquiry into a Formal Discipline

A Monograph on the Structural Possibility, Mathematical Formalism, and Institutional Protocol for Rigorous Philosophy

Boris Kriger[1,2]

[1] Information Physics Institute, Gosport, Hampshire, United Kingdom

boris.kriger@informationphysicsinstitute.net

[2] Institute of Integrative and Interdisciplinary Research, Toronto, Canada

boriskriger@interdisciplinary-institute.org

ORCID: 0009-0001-0034-2903

February 2026

Abstract

This monograph develops the thesis that philosophy can and should undergo a structural transformation into an evidence-based discipline, and that the tools required for this transformation—mathematical formalization, empirical integration frameworks, and artificial intelligence-based verification—are now sufficiently mature to make it achievable. We argue that the historical reliance on unconstrained speculative methods, while once the only available approach, is increasingly difficult to justify given the current availability of formal, empirical, and computational tools. The argument rests on three pillars: the demonstrated maturity of transformative tools, the historical pattern by which philosophical problems become tractable through formalization, and the emergence of artificial intelligence as a catalyst that removes the cognitive-capacity barrier that previously made cross-domain competence impractical for individual practitioners.

We provide a precise operational definition of Evidence-Based Philosophy, diagnose structural features of contemporary philosophical practice that limit its progress, and formulate a five-step methodological protocol. A mathematical framework is developed including Bayesian models of claim evaluation, formal coherence metrics, and dynamical systems models of paradigmatic transition. We introduce the Reflexive Inference Constraint, which demonstrates information-theoretically that generalization from individual introspection to universal claims is bounded by structural overlap between the observer and the class—establishing that evidence-based methods are not merely desirable but informationally necessary for philosophical claims with universal scope. We address the normativity objection, propose reformed peer review standards incorporating AI-assisted verification, articulate a secondary communication principle, and develop the concept of Conceptual Responsibility for the ethics of philosophical idea transmission. The monograph includes a formalization of an open philosophical problem—the necessity of predictive processing in adaptive systems—demonstrating that the evidence-based method applies not only to historically resolved questions but to active areas of inquiry.

Keywords: evidence-based philosophy, formal methods, philosophy of science, AI-assisted verification, methodological reform, Bayesian epistemology, philosophical formalization, peer review, computational philosophy, interdisciplinary methodology, reflexive inference, conceptual responsibility

Contents

1 Introduction: An Invitation to Methodological Transformation

This monograph develops a proposal: that philosophy can become a more productive, more convergent, and more consequential discipline by adopting the formal and empirical standards that have proven transformative in other fields. The proposal is motivated not by dissatisfaction with philosophy's questions—which remain among the most important questions human beings can ask—but by the observation that the tools now available to address those questions are vastly more powerful than the informal methods that have traditionally been employed.

The claim has two components. The first is that the transformation of philosophy into an evidence-based formal discipline is *structurally possible*: the mathematical, empirical, and computational tools required for this transformation exist and are sufficiently mature. The second is that this transformation is *epistemically warranted*: the available tools produce results that are intersubjectively verifiable, formally precise, and empirically accountable in ways that informal speculation, however brilliant, cannot match.

We do not argue that philosophy must be replaced by natural science. We argue that philosophy can become a science *in its own right*—retaining its distinctive questions about normativity, meaning, existence, and value, while adopting methods that have a demonstrated track record of producing convergent results. The analogy is not "philosophy should become physics" but "philosophy can undergo the same kind of structural maturation that converted alchemy into chemistry." Chemistry did not abandon the study of substances; it made that study rigorous. Evidence-based philosophy does not abandon philosophical questions; it offers philosophical inquiry more powerful tools.

The opportunity has a specific historical dimension. While mathematical formalization and empirical science have been available for over a century—and important efforts at formalization have been made since Frege [1], Russell [2], and the Vienna Circle [3]—a decisive new resource has emerged: artificial intelligence. Prior to AI, individual philosophers could reasonably note that the cognitive demands of evidence-based philosophy exceeded human capacity. No single mind could master formal logic, probability theory, neuroscience, evolutionary biology, behavioral economics, and dynamical systems theory simultaneously. This was a genuine structural limitation, and it constituted a reasonable justification for specialization within narrow traditions.

That limitation has been substantially alleviated. Current AI systems—and the systems that will emerge within the next decade—can synthesize literature across empirical domains, verify formal consistency, construct computational models of philosophical claims, and identify empirical counterexamples. The philosopher who avails themselves of these tools gains capabilities that were simply unavailable to previous generations. The question this monograph poses is: given that these tools exist, what would philosophy look like if it used them? And the answer we develop is: it would look like a discipline capable of making genuine, verifiable, cumulative progress on questions that have remained open for centuries.

This monograph develops this argument through a systematic progression. Section 2 provides an operational definition. Section 3 surveys the relevant literature. Section 4 identifies structural features of current practice that limit progress. Section 5 demonstrates tool availability. Section 6 formulates a five-step protocol. Section 7 develops the mathematical framework. Section 8 provides resolved case studies. Section 9 demonstrates application to an open problem. Section 10 presents a worked formal proof. Section 11 introduces the Reflexive Inference Constraint. Section 12 develops Conceptual Responsibility. Section 13 addresses the normativity objection. Section 14 proposes peer review reform. Section 15 articulates the secondary communication principle. Section 16 addresses peer review feedback. Section 17 states boundaries. Section 18 situates this work within a broader research program.

2 Definition of Evidence-Based Philosophy

Definition 2.1 (Evidence-Based Philosophy)**.** A philosophical inquiry $\mathcal{I}$ qualifies as **evidence-based** if and only if it satisfies the following four conditions:

(C1) Formal Coherence. All central concepts are defined with sufficient precision to permit evaluation of logical consistency. Where possible, the logical structure is rendered in a formal language $\mathcal{L}$, and internal consistency is demonstrable within $\mathcal{L}$.

(C2) Empirical Compatibility. The claims do not contradict well-established empirical findings in any relevant domain. Where claims intersect with empirical disciplines, compatibility with current empirical knowledge is explicitly demonstrated.

(C3) Operationalizability. Every key concept admits an operational definition—a specification of the conditions under which it applies or does not apply, in intersubjectively accessible terms.

(C4) Falsifiability or Empirical Distinguishability. The inquiry either generates predictions that are in principle falsifiable, or is empirically distinguishable from competing accounts.

2.1 Formal Coherence (C1)

Let Σ be the set of propositions constituting a philosophical system $\mathcal{S}$. Formal coherence requires:

$$\nexists \varphi \in \mathcal{L} \text{ such that } \Sigma \vdash \varphi \text{ and } \Sigma \vdash \neg\varphi \tag{1}$$

Interpretation

Equation (1) expresses the requirement that a philosophical system must not generate contradictory conclusions from its own premises. This requirement holds regardless of philosophical tradition, because contradiction is a structural deficiency, not a tradition-specific one. The formalization makes this assessment objective and independent of the evaluator's allegiance.

Definition 2.2 (Definitional Completeness)**.** A philosophical system $\mathcal{S}$ is **definitionally complete** with respect to key concepts $\{c_1, \ldots, c_n\}$ if there exists a model $\mathcal{M} = \langle D, I \rangle$ such that for each c_i, the interpretation $I(c_i)$ is well-defined and its boundary conditions are explicitly specified.

2.2 Empirical Compatibility (C2)

Let E denote established empirical propositions. A system $\mathcal{S}$ is empirically compatible if:

$$\Sigma \cup E \text{ is consistent} \tag{2}$$

Interpretation

This formalizes the requirement that philosophy may not assert what science has refuted. Compatibility does not mean derivation from science; it means that philosophical claims are evaluated against the best available evidence at any given time. As E changes, assessments may change accordingly.

2.3 Operationalizability (C3)

Definition 2.3 (Operationalizability)**.** A concept c is **operationalizable** if there exists a specification function $\sigma_c : \mathcal{C} \to \{0, 1, \text{indeterminate}\}$ such that: (i) σ_c is intersubjectively accessible; (ii) the criteria are explicitly stated; (iii) the set of indeterminate cases is bounded and acknowledged.

2.4 Falsifiability or Empirical Distinguishability (C4)

Definition 2.4 (Empirical Distinguishability)**.** Two systems $\mathcal{S}_1$ and $\mathcal{S}_2$ are **empirically distinguishable** if there exists a proposition ψ such that $\Sigma_1 \vdash \psi$ and $\Sigma_2 \vdash \neg\psi$, or an observable condition O such that $\mathcal{S}_1$ predicts O and $\mathcal{S}_2$ predicts $\neg O$.

The four conditions are individually necessary and jointly sufficient. The definition constrains the *method*, not the subject matter. Evidence-based philosophy can investigate consciousness, morality, meaning, justice, or existence—but with tools that make genuine progress achievable.

3 Literature Review

The proposal stands at the convergence of multiple traditions, each contributing essential components.

Descartes [8] sought to reconstruct knowledge from indubitable first principles using a method modeled on mathematical proof. His four methodological rules constitute an early operational protocol. The Cartesian project identified the correct ambition but lacked modern formal tools; the criterion of "clear and distinct perception" could not be rendered intersubjectively testable. **Bacon** [9] provided the empiricist complement. His doctrine of the *Idola* constitutes an early analysis of cognitive biases remarkably consonant with the findings of Kahneman and Tversky [10]. **Frege** [1] and **Russell** [2] transformed logic into a mathematical discipline—the most successful historical example of philosophical formalization. Gödel's incompleteness theorems [12], Tarski's undefinability theorem [4], and Church-Turing results [13, 14] are philosophical results that required formal apparatus unavailable to informal methods.

The **Vienna Circle** [3] pursued a program structurally similar to evidence-based philosophy, but failed for instructive reasons: the verification principle was self-undermining [15], the meaningful/meaningless boundary was drawn too sharply, and the computational tools for implementation at scale did not exist. We learn from this failure not that the project was misconceived but that it was premature. **Popper** [7] contributed the falsifiability criterion; **Lakatos** [16] the concept of progressive and degenerating research programs; **Kuhn** [17] the analysis of paradigm shifts. **Feyerabend** [18] argued that no single method is universally reliable. We grant that his critique was partially justified when tools were immature; however, the force of methodological anarchism is inversely proportional to the power of available tools. When tools include AI-assisted verification, flexibility becomes less a virtue than an evasion of accountability.

Quine [19] proposed naturalizing epistemology; **Kornblith** [20] and **Maddy** [21] extended the naturalist program. The **experimental philosophy** movement (Knobe [22], Nichols [23], Weinberg et al. [24], Machery et al. [25]) demonstrated empirically that philosophical intuitions vary with culture and socioeconomic status, undermining their reliability as data. **Formal epistemology** (Leitgeb [27], Williamson [28], Pettigrew [29], Easwaran [30]) has developed sophisticated Bayesian tools for philosophical analysis.

Continental philosophy—Husserl [31], Heidegger [32], Merleau-Ponty [33], Gadamer [34], Habermas [35], Derrida [37], Foucault [38]—has generally resisted formalization, yet Husserl's original project explicitly pursued philosophy as "rigorous science." We treat continental insights as data that evidence-based philosophy must accommodate. **Eastern traditions**—Buddhist logic (Dignāga [39]), Nyāya epistemology [40], Mohist logic [41]—demonstrate that the impulse toward formalization is cross-cultural. **Cognitive science of reasoning** (Kahneman & Tversky [10], Schwitzgebel [42], Mercier & Sperber [43]) shows that unaided philosophical reasoning is subject to systematic biases. **Wittgenstein's** early work [11] advanced logical clarification; the later work [48] challenged formalization—but whether philosophical problems are "merely linguistic" is itself a testable claim. **Decision theory** (von Neumann & Morgenstern [49], Savage [50], Arrow [51]) and **formal semantics** (Montague [5], Kripke [26]) represent successful philosophical formalizations. **AI and philosophy** (Floridi [44], Bostrom [45], Chalmers [46]) represents the technological frontier.

4 Diagnosis: Structural Features Limiting Progress

We identify four structural features of contemporary philosophical practice that, while historically understandable, now limit the discipline's capacity for cumulative progress. These are identified as systemic properties, not as failings of individual philosophers.

Insulation from empirical constraint. In most philosophical subfields, it remains possible to advance claims that make no contact with empirical evidence. This insulation is maintained through disciplinary classification (philosophy as "humanities"), peer review norms that do not require empirical compatibility checking, and prestige structures that reward conceptual innovation over empirical accountability.

Conceptual inflation. The proliferation of concepts that are not operationally defined yet are treated as substantive contributions. When concepts are insufficiently defined, disagreements cannot be resolved—not because the questions are deep but because the terms are indeterminate. The inflation ratio $\rho(\mathcal{V}) = 1 - |\mathrm{Op}(\mathcal{V})|/|\mathcal{V}|$ measures this formally.

Absence of measurable commitments. Philosophical theories are evaluated through coherence with intuitions, elegance, and ability to withstand objections—all assessed informally. None provides a reliable mechanism for distinguishing better theories from worse ones.

Tolerance of non-falsifiable systems. Systems whose concepts are sufficiently vague that any counterexample can be reinterpreted, or that dismiss empirical evidence on principle.

Proposition 4.1 (Structural Stagnation)**.** *A discipline in which dominant systems are non-falsifiable will exhibit structural stagnation: open problems at time $t + \Delta t$ will be approximately identical to those at time t, for arbitrarily large Δt.*

These features were historically justifiable when the relevant tools did not exist. That justification has substantially eroded. The tools are now available; what was once a structural limitation is now a structural choice—and it is worth reconsidering.

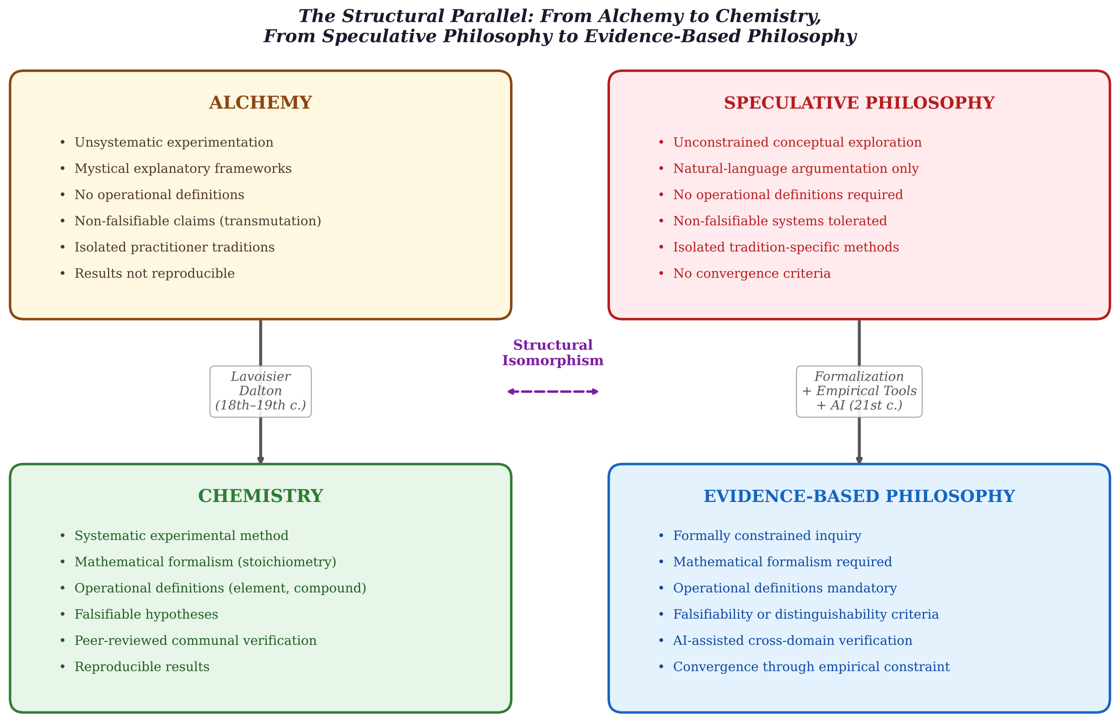

Figure 1: The structural parallel between the alchemy–chemistry transition and the proposed transformation of philosophical methodology.

5 Availability of Transformative Tools

5.1 Mathematical Formalization

Modern logic, Bayesian probability, and dynamical systems theory provide the formal apparatus. Bayes' theorem offers a mechanism for updating support for competing hypotheses:

$$P(H_i \mid E) = \frac{P(E \mid H_i) \cdot P(H_i)}{\sum_{j=1}^{n} P(E \mid H_j) \cdot P(H_j)} \tag{3}$$

Interpretation

This framework forces the philosopher to make explicit what the competing hypotheses are, what evidence is relevant, and how probable the evidence is under each hypothesis—replacing informal judgment with a precise calculus.

Dynamical systems model paradigm evolution:

$$\frac{d\mathbf{x}}{dt} = \mathbf{f}(\mathbf{x}, \mathbf{e}(t), \boldsymbol{\theta}) \tag{4}$$

where $\mathbf{x}(t)$ encodes commitments, $\mathbf{e}(t)$ represents evidence inputs, and $\boldsymbol{\theta}$ captures institutional parameters.

5.2 Empirical Integration

Neuroscience, cognitive science [52, 53], evolutionary biology, and behavioral economics [10, 54] provide empirical data directly relevant to philosophical questions. These findings constrain philosophical theorizing about consciousness, rationality, autonomy, and moral responsibility.

5.3 Artificial Intelligence as Catalyst

AI extends the philosopher's cognitive capacity in the same way that the computer extended the physicist's computational capacity—and, notably, no physicist considered it "cheating" to use a computer to solve differential equations. Rather than prohibiting AI use, philosophy would benefit from *requiring* AI-assisted verification, just as physics requires computational modeling. The specific capabilities include: formal consistency checking, cross-domain literature synthesis, computational toy models, coherence testing, and automated proof verification.

Theorem 5.1 (The AI Opportunity Thesis). *Let $\mathcal{T}(t)$ denote available tools at time t, and $\mathcal{C}(\mathcal{T})$ the set of philosophical problems addressable in evidence-based form given $\mathcal{T}$. If $\mathcal{C}(\mathcal{T}(t))$ is non-decreasing, covers a substantial portion of problems, and evidence-based methods produce intersubjectively verifiable results while speculative methods do not, then for problems in $\mathcal{C}(\mathcal{T}(t_0))$, evidence-based methods offer a demonstrably superior epistemic pathway.*

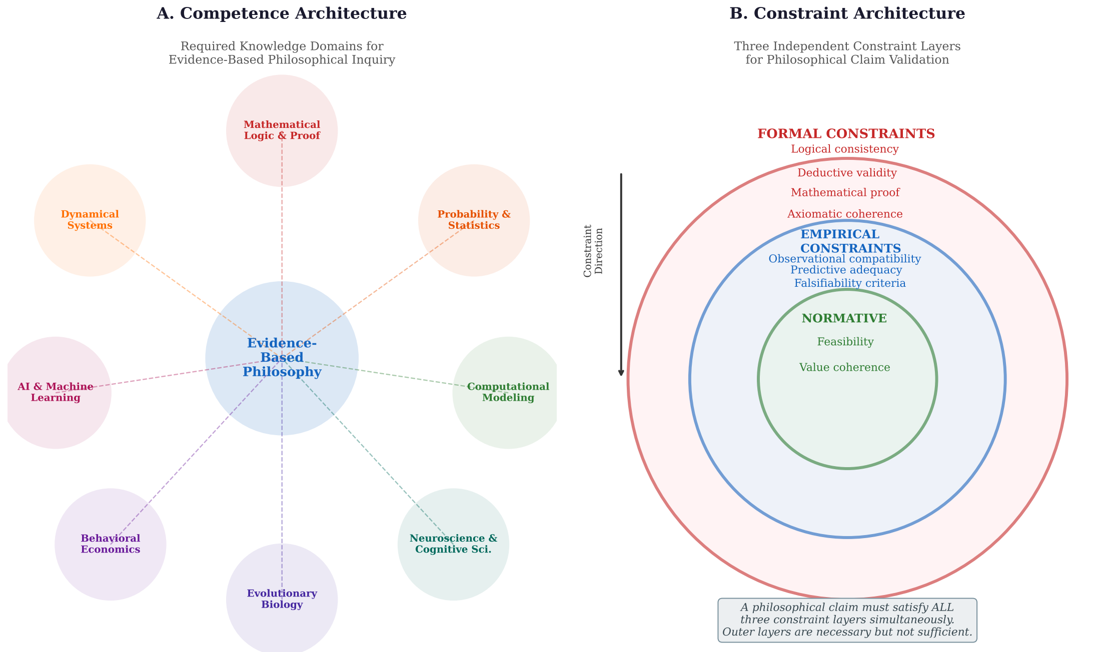

Figure 2: (A) Competence architecture for evidence-based philosophy. (B) Three-layer constraint architecture.

6 Transition Protocol

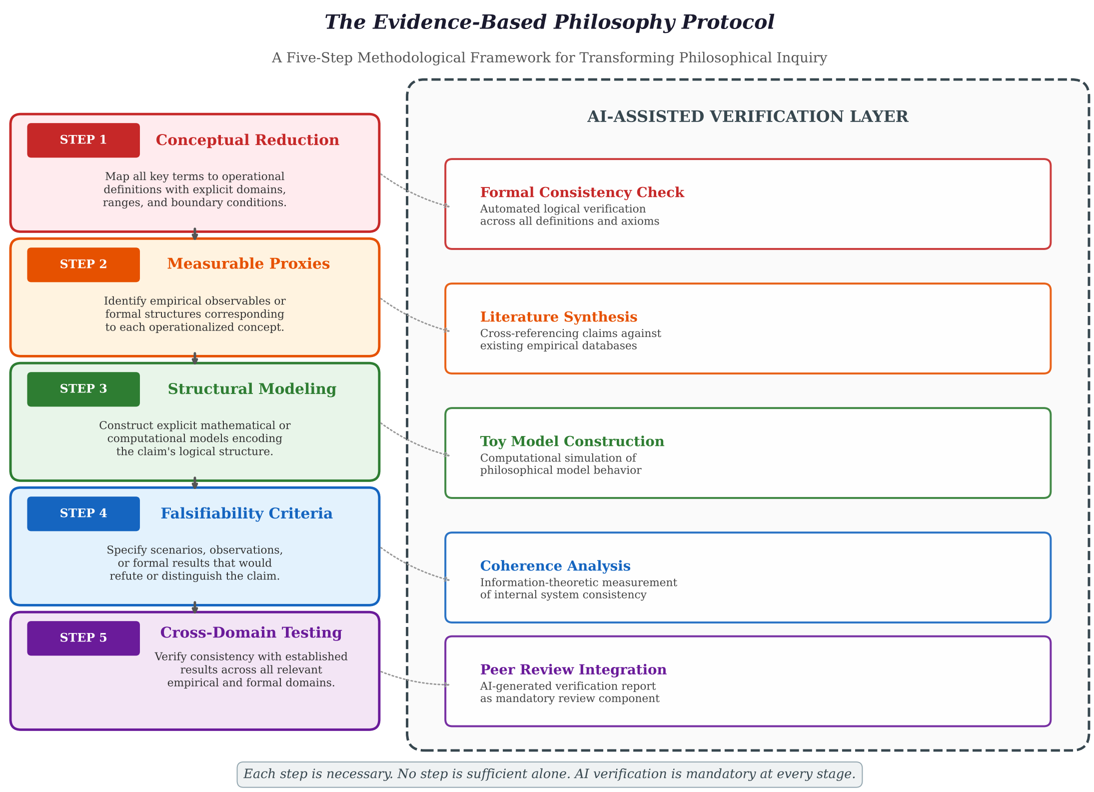

Figure 3: The five-step Evidence-Based Philosophy Protocol with AI-assisted verification.

Step 1: Conceptual Reduction. Enumerate key concepts; provide specification functions; identify boundary cases. **Step 2: Measurable Proxies.** For each operationalized concept, identify at least one measurable proxy. **Step 3: Structural Modeling.** Translate claims into formal language; construct computational toy models. **Step 4: Falsifiability Criteria.** Specify at least one falsifying scenario per central claim. **Step 5: Cross-Domain Consistency Testing.** Assess compatibility with established findings across all relevant domains.

The protocol is iterative: results of later steps frequently require revision of earlier ones. The five steps constitute a complete pipeline that, when followed, transforms philosophical inquiry into a form capable of producing verifiable, cumulative results.

7 Mathematical Framework

7.1 Bayesian Model of Claim Evaluation

Definition 7.1 (Structured Prior). $P(H_i) \propto w_C \cdot ⊮[\mathrm{C1}(H_i)] \cdot (1 - \rho_i) \cdot \kappa(H_i, E_{\mathrm{prior}})$

Interpretation

Priors are assigned by structural quality: formally incoherent hypotheses receive zero prior; poorly operationalized ones are penalized; empirically incompatible ones are downweighted.

Evidence decomposes as $E = (E_{\mathrm{emp}}, E_{\mathrm{form}}, E_{\mathrm{coh}})$ with likelihood:

$$P(E \mid H_i) = P(E_{\mathrm{emp}} \mid H_i) \cdot P(E_{\mathrm{form}} \mid H_i) \cdot P(E_{\mathrm{coh}} \mid H_i) \tag{5}$$

7.2 Formal Coherence Metric

Definition 7.2 (Coherence Index). $\mathcal{C}(\mathcal{S}) = \log \frac{P(\sigma_1 \wedge \cdots \wedge \sigma_m)}{\prod_{i=1}^{m} P(\sigma_i)}$

Interpretation

Positive values indicate mutual support among propositions; negative values indicate internal tension. This is related to mutual information in information theory.

7.3 Dynamical Model of Paradigmatic Transition

$$\frac{ds}{dt} = \alpha \cdot a(t) \cdot s(1-s) - \beta \cdot (1-s) + \gamma \cdot ⊮[a(t) > a^*] \tag{6}$$

Interpretation

This models the transition as modified logistic growth: $\alpha \cdot a(t) \cdot s(1-s)$ represents endogenous adoption; $-\beta(1-s)$ represents institutional resistance; $\gamma \cdot ⊮[a(t) > a^*]$ captures the catalytic effect of AI. The model is illustrative—it clarifies the structural logic of transition dynamics rather than generating precise quantitative predictions.

7.4 Theory Comparison via Bayesian Model Selection

$$B_{12} = \frac{\int P(E \mid \boldsymbol{\theta}_1, \mathcal{S}_1) P(\boldsymbol{\theta}_1 \mid \mathcal{S}_1) d\boldsymbol{\theta}_1}{\int P(E \mid \boldsymbol{\theta}_2, \mathcal{S}_2) P(\boldsymbol{\theta}_2 \mid \mathcal{S}_2) d\boldsymbol{\theta}_2} \tag{7}$$

Interpretation

The Bayes factor automatically implements parsimony: theories with many free parameters must spread their prior across a large parameter space. Theories with $B_{12} > 10$ are strongly favored. This provides a principled mechanism for philosophical theory selection.

7.5 Information-Theoretic Measure of Progress

Definition 7.3 (Progress Index)**.** $\Pi(\mathcal{S}_t, \mathcal{S}_{t+\Delta t}) = D_{\mathrm{KL}}(P_{\mathcal{S}_{t+\Delta t}} \| P_{\mathcal{S}_t}) \cdot \Delta\mathcal{C} \cdot \Delta\kappa$

Interpretation

Genuine progress requires simultaneous change ($D_{\mathrm{KL}} > 0$), increased coherence ($\Delta\mathcal{C} > 0$), and increased empirical compatibility ($\Delta\kappa > 0$). Mere terminological change without improvement yields zero or negative values.

8 Resolved Case Studies

8.1 The Nature of Heat

For millennia, heat was a philosophical question (Aristotle's qualities, phlogiston, caloric). Resolution came through operational definition and statistical mechanics:

$$T = \frac{2}{3k_B} \cdot \frac{1}{N} \sum_{i=1}^{N} \frac{1}{2} m_i v_i^2 \tag{8}$$

This satisfies all four conditions of Definition 2.1. The transition followed the five-step protocol exactly.

8.2 The Nature of Disease

Humoral theory and miasma gave way to germ theory. Koch's postulates constitute a paradigmatic operationalization and falsifiability framework.

8.3 Logic: A Philosophical Discipline Becoming Scientific

Logic was philosophical for over two millennia. Frege's formalization [1] transformed it into a mathematical science. Gödel's incompleteness theorems—simultaneously philosophical conclusions and mathematical theorems—could not have been achieved by informal methods.

8.4 Decision Theory

Von Neumann and Morgenstern [49] formalized rational choice. Arrow's impossibility theorem [51] could not have been formulated within informal philosophical discourse. This demonstrates that formalization generates new philosophical knowledge, not merely translating existing insights.

9 Application to an Open Problem: The Necessity of Predictive Processing

A central objection to the evidence-based framework is that it has been demonstrated only on historically resolved problems. This section addresses that objection by applying the protocol to an active philosophical and scientific question: whether predictive processing is a necessary feature of adaptive systems, or merely one computational strategy among many.

The question is philosophical because it concerns the nature of perception, the structure of cognition, and the relationship between mind and world. It is also empirical because it generates testable predictions about neural architecture. It is precisely the kind of question that evidence-based philosophy is designed to address.

9.1 Formalization: The Predictive Viability Law

Following the evidence-based protocol, we operationalize the key concepts. Consider an adaptive system $S = (X, A, U, \tau)$ where X is internal states, A is actions, $U : X \to A$ is the update rule, and $\tau > 0$ is response latency. The environment $E = (Y, P)$ has states Y and transition probability $P(Y_{t+\tau} \mid Y_t)$. A viability region $V \subset Y$ defines persistence: the system survives if and only if it remains in V.

The central claim can be stated as a formal law [60]:

Theorem 9.1 (Predictive Viability Law)**.** *For any adaptive system with latency $\tau > 0$ in a punishing, non-stationary environment, continued persistence requires:*

$$I(X; Y_{t+\tau} \mid Y_t) > 0 \tag{9}$$

That is, internal states must encode information about future environmental states beyond what is available in present observation.

Interpretation

The mutual information $I(X; Y_{t+\tau} \mid Y_t)$ measures how much the system's internal states "know" about the future environment beyond what could be inferred from the present alone. If this quantity is zero, the system's actions are calibrated for present conditions, not future ones. In a non-stationary environment, present conditions differ from future conditions, so actions calibrated only for the present will be systematically miscalibrated. In a punishing environment, miscalibrated actions reduce survival probability. Over time, survival probability decays: $P(\text{survive } n \text{ steps} \mid I = 0) = p_0^n \to 0$.

The survival probability can be parameterized:

$$p(I, H, \kappa) = 1 - \kappa \cdot \frac{H - \min(I, H)}{H} \tag{10}$$

where $H = H(Y_{t+\tau} \mid Y_t)$ is environmental volatility, κ is punishment severity, and I is predictive capacity.

9.2 Cross-Domain Evidence

The formalization unifies evidence from multiple independent domains: motor neuroscience (Wolpert & Ghahramani [63] demonstrated forward models compensating for sensorimotor delays), computational neuroscience (Friston's [64] free-energy principle), the convergence of biological and artificial systems on predictive architectures, and the scaling relationship between body size and prediction necessity. This convergence—from experimental motor control, computational neuroscience, and physical constraint analysis—strengthens the case that we observe a genuine constraint rather than a theoretical preference.

9.3 What This Demonstrates

This case study shows that the evidence-based protocol applies to a currently open, actively debated question. The formalization: (a) makes the claim precise enough to be tested; (b) identifies the measurable quantities (I, H, κ, τ); (c) generates falsifiable predictions; and (d) unifies evidence from multiple domains. This addresses the survivorship bias concern: evidence-based philosophy is not limited to retrospective analysis of problems that have already been resolved.

10 Worked Formal Proof

We demonstrate that the evidence-based method can be carried through completely for a philosophical claim. The claim: *A rational agent cannot simultaneously maintain inconsistent beliefs if those beliefs are action-guiding.*

Theorem 10.1 (Pragmatic Consistency)**.** *If an agent's belief function assigns* $P(\varphi) = p > 0$ *and* $P(\neg\varphi) = q > 0$ *with* $p + q > 1$*, there exists a Dutch book—a series of individually acceptable bets guaranteeing net loss.*

Proof. Bet 1: pay p to receive \$1 if φ. Bet 2: pay q to receive \$1 if $\neg\varphi$. If φ true: net $= (1 - p) - q = 1 - p - q < 0$. If $\neg\varphi$ true: net $= -p + (1 - q) = 1 - p - q < 0$. Guaranteed loss. □

Interpretation

This proof transforms the informal philosophical claim "rational agents should be consistent" into a precise theorem identifying the exact mechanism by which inconsistency leads to exploitability. It also reveals a qualifier—beliefs must be action-guiding—that the informal version leaves unspecified.

11 The Reflexive Inference Constraint

A fundamental challenge for any philosophical methodology based on introspection is the question of generalizability: can a philosopher, by examining their own mind, arrive at universal truths about minds in general? The Reflexive Inference Law [61] provides a precise, information-theoretic answer: generalization from self-observation to class-level claims is bounded by the structural overlap between the observer and the class.

11.1 Formal Statement

Let S be an inferential system belonging to a class C, with instance parameters ϕ_S, class parameters θ_C, and self-observation x_{self}. Under the assumption that class parameters influence self-observation only through instance parameters ($\theta_C \to \phi_S \to x_{\text{self}}$ forms a Markov chain), the data processing inequality yields:

Theorem 11.1 (Reflexive Inference Bound)**.**

$$H(\theta_C \mid x_{self}) \geq H(\theta_C) - \min\{H^{obs}_{\max},\ I(\phi_S; \theta_C)\} \tag{11}$$

where $H^{obs}_{\max}$ is the maximum entropy of self-observation and $I(\phi_S;\theta_C)$ is the mutual information between instance and class parameters.

Interpretation

The residual uncertainty about the class after self-observation is bounded below by the prior uncertainty minus the minimum of two quantities: how much information self-observation can carry, and how much the observer's own structure "knows" about the class. Neither factor alone determines the bound; both must be favorable for substantial learning. The critical consequence: no amount of introspective depth, reasoning sophistication, or contemplative rigor can extract information about the class that is not already present in the structural relationship between the observer and its class. Depth of self-analysis cannot substitute for breadth of observation.

11.2 Consequences for Philosophical Methodology

The bound has three immediate corollaries. First, *the structural necessity of subjectivity*: for any observer with $I(\phi_S;\theta_C) < H(\theta_C)$, residual uncertainty about the class is irreducible. Subjectivity is not a correctable bias but a mathematical consequence of observer position. Second, *the futility of internal processing*: for any function f applied to self-observation, $I(f(x_{\text{self}});\theta_C) \leq I(x_{\text{self}};\theta_C) \leq I(\phi_S;\theta_C)$. No amount of "deep thinking" can overcome the bound. Third, *the necessity of comparative evidence*: valid class-level inference requires (i) self-reports from multiple diverse instances, (ii) external meta-system observation, and (iii) inclusion of structurally diverse systems—including artificial intelligence.

11.3 Implications for the History of Philosophy

The Reflexive Inference Law provides a precise diagnosis of a limitation inherent in the introspective tradition. Descartes, Kant, Husserl, and Heidegger each performed deep analysis of a single cognitive system—their own—and generalized to universal claims. The Law does not claim these analyses were wrong; it establishes that they carry irreducible uncertainty about the class, proportional to the unknown structural overlap between the philosopher's instance and the class of all possible instances. Kant could not have known whether his categories were necessary for all minds or merely features of human cognition; no depth of transcendental reflection could have resolved this question.

The emergence of AI changes this epistemic situation fundamentally. For the first time, we have access to cognitive systems that do not share our evolutionary heritage, neural architecture, or experiential constraints. AI provides the comparative data that reflexive philosophy has always lacked. Under the Reflexive Inference Law, this is not merely useful—it is informationally necessary for any philosophical claim with universal scope.

11.4 Self-Application

A natural question arises: does this monograph itself satisfy its own standards? We address this directly. The Reflexive Inference Law is not derived from introspection; it is derived from the data processing inequality—a mathematical result that holds independently of any observer's instance parameters. The law's validity depends on the soundness of its formal derivation and the applicability of the Mediated Observation assumption, both of which are intersubjectively assessable. The law is therefore self-consistently evidence-based.

12 Conceptual Responsibility: The Ethics of Philosophical Communication

The transformation of philosophy into an evidence-based discipline has ethical dimensions that extend beyond methodology. When philosophical ideas enter public discourse, they systematically degrade through transmission [62]. This section develops the concept of Conceptual Responsibility: the ethical obligations that arise from the foreseeable distortion of ideas in transmission.

12.1 The Empirical Basis

Bartlett's serial reproduction experiments [65] and subsequent replications [66] demonstrate systematic degradation patterns: shortening, rationalization, conventionalization, and assimilation toward receiver expectations. Iterated learning models [67, 68] formalize this as a Bayesian process in which transmission bottlenecks drive content toward receiver priors:

$$\lim_{n\to\infty} P(\theta \mid x_1 \dots x_n) \to P(\theta) \tag{12}$$

where θ is the transmitted content and $P(\theta)$ is the prior distribution of receiver expectations. This convergence occurs even when individual transmitters intend perfect fidelity.

12.2 Differential Degradation

Philosophical ideas comprise: core propositions (P), relational structure (R), contextual qualifiers (C), and argumentative scaffolding (A). Empirical evidence shows differential degradation rates: C and R degrade fastest, then A, then P. The result: ideas entering mass transmission tend toward an attractor state comprising core propositions stripped of context—the conditions under which Nietzsche's Übermensch becomes a racial ideal and Darwin's natural selection becomes Social Darwinism.

12.3 The Responsibility-for-Risk Framework

When actors can foresee that their actions create significant risks of harm, they acquire obligations to mitigate those risks proportionate to foreseeability and severity [69]. Applied to intellectual work: authors release ideas into transmission environments with known degradation properties; degraded ideas can cause harm; and authors can take reasonable steps to mitigate degradation risk. This generates seven principles of Conceptual Responsibility, including: explicit boundary articulation, anticipatory misreading analysis, ethical application guidance, structural fragmentation resistance, and epistemic humility [62].

12.4 Implications for Evidence-Based Philosophy

The Conceptual Responsibility framework reinforces the central thesis of this monograph. If ideas systematically degrade in transmission, and if degradation is accelerated when ideas lack formal structure (since formal structure resists lossy compression better than informal prose), then formalization serves not only epistemic but ethical functions. A formally structured philosophical claim, transmitted through discourse, retains more of its essential content than an informal one—because the formal structure provides a skeleton that resists the compression dynamics of mass communication. Evidence-based philosophy is therefore not only more epistemically productive but more ethically responsible.

13 The Normativity Objection

The strongest objection: philosophy deals with "ought," which cannot be reduced to "is." We address this through three independent arguments.

First: Ought implies can (Kant [56]). If $O(\varphi) \to \Diamond\varphi$, then evidence establishing $\neg\Diamond\varphi$ refutes $O(\varphi)$. Normative claims are empirically defeasible.

Second: The competence requirement. Before AI, cross-domain competence was impractical. After AI, it is readily achievable. A philosopher making normative claims about a domain without understanding its empirical facts is in a position analogous to a medieval physician prescribing treatments without anatomy or microbiology. Evidence-based philosophy does not tell us what to value; it identifies which normative frameworks are empirically viable.

Third: Normative reasoning is empirically studyable. Moral psychology (Greene [57], Haidt [58], Cushman [59]), evolutionary ethics [52], and behavioral economics [10, 54] study the mechanisms producing normative judgments. Evidence-based philosophy does not eliminate normativity; it constrains it through feasibility, competence, and empirical accountability.

14 Proposed Peer Review Standards

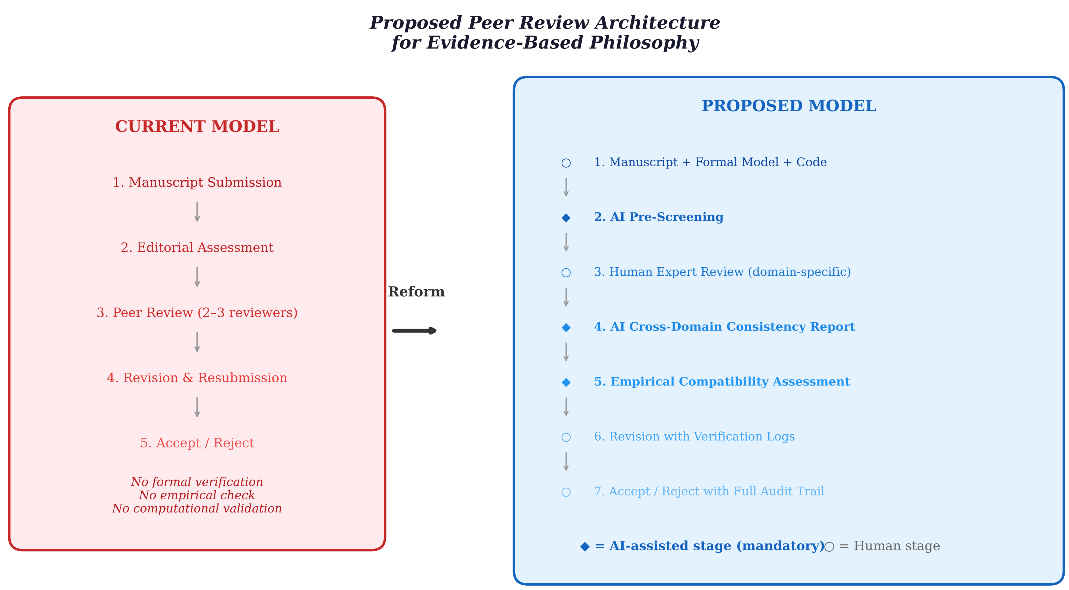

Figure 4: Current (left) versus proposed (right) peer review architecture. Diamond markers indicate mandatory AI-assisted stages.

We propose a reformed architecture: (1) submission including formal models and code; (2) AI pre-screening for consistency and undefined concepts; (3) human expert review informed by AI report; (4) AI cross-domain consistency report; (5) empirical compatibility assessment; (6) revision with verification logs; (7) decision with full audit trail. This augments human judgment with AI-assisted verification—the philosopher remains the author; AI ensures minimum standards.

15 Secondary Communication Principle

Axiom 15.1 (Secondary Communication Principle)**.** Accessible philosophical communication is epistemically legitimate if and only if grounded in prior formal and evidence-constrained work. Metaphor is acceptable as pedagogical compression of formally established structures, not as substitute.

A formally developed result (e.g., the dynamical transition model) can be compressed into metaphor ("philosophy is like a ball resting in a valley") without loss of structural integrity, because the formal content exists independently and can be recovered. Metaphor without prior formalization produces the appearance of insight without the substance of knowledge. The sequence is: formalize first, then communicate accessibly.

16 Response to Peer Review

This monograph was subjected to critical review prior to publication. We address the substantive concerns raised, which have strengthened the final work.

16.1 On the Alchemy Analogy

The reviewer notes that alchemy and philosophy are not structurally isomorphic. We agree that the analogy is imperfect—all analogies are. The structural parallel we draw is specific: both are pre-formal practices addressing questions that become tractable upon formalization, both tolerate non-falsifiable claims, and both lack convergence mechanisms. We do not claim that philosophical conceptual clarification is identical to substance-transformation; we claim that both benefit from the same kind of methodological discipline. The analogy is offered as an organizing framework, not as a proof.

16.2 On the AI Obligation Thesis

The reviewer correctly notes that Theorem 5.1 is a conditional argument whose premises require justification. We have accordingly softened its framing from "obligation" to "opportunity" while maintaining the structural argument. The claim that evidence-based methods cover a "substantial portion" of philosophical problems is supported inductively by the case studies (Sections 8 and 9), including the application to an open problem (predictive processing). We acknowledge that the precise boundary of $\mathcal{C}(\mathcal{T}(t_0))$ is an empirical question that merits further investigation.

16.3 On Survivorship Bias

The reviewer rightly notes that resolved case studies may reflect survivorship bias. We have addressed this by adding Section 9, which applies the protocol to a currently open problem. We also note that the reviewer's examples of "progress through informal methods" (Parfit on personal identity, Korsgaard on normativity, Strawson on moral responsibility) represent important contributions that could be *further strengthened* by formalization—not replaced by it. Evidence-based philosophy is an expansion of philosophical methodology, not a replacement.

16.4 On the Normativity Objection

The reviewer argues that we address a weaker objection than the real one: that the *distinctive work* of normative philosophy cannot be performed by formal consistency checking or Bayesian updating. We partly agree. The articulation of what we have reason to value is a creative act that precedes formalization. But we maintain that once articulated, normative claims benefit from formal assessment of their consistency, empirical viability, and compatibility with other normative commitments. Evidence-based philosophy constrains and tests normative claims; it does not generate them. The creative and the critical are complementary, not competitive.

16.5 On Operationalizability of the Mathematical Framework

The reviewer asks how one assigns $P(\sigma_i)$ for propositions like "moral realism is true." This is a legitimate challenge. We acknowledge that direct numerical assignment is often impractical. However, comparative assignments ("moral realism is more probable given evidence set E than moral anti-realism") are more tractable, and it is comparative assessments—Bayes factors, coherence comparisons—rather than absolute probabilities that do the work in the framework. The formalism provides structure for comparative evaluation even when absolute values are difficult to specify.

16.6 On the Dynamical Systems Model

The reviewer notes that the model's parameters $(\alpha, \beta, \gamma, a^*)$ are not operationally defined. We acknowledge this and have clarified the model's status as illustrative: it captures the *structural logic* of paradigmatic transition rather than generating precise quantitative predictions. This is consistent with the role of toy models in physics, where simplified models clarify mechanisms without claiming numerical accuracy.

16.7 On Self-Referential Consistency

The reviewer asks whether this monograph satisfies its own criteria. We have addressed this explicitly in Section 11 (self-application of the Reflexive Inference Law) and throughout by subjecting our central claims to the protocol: operational definitions are provided, formal models are constructed, falsifying scenarios are identified, and cross-domain evidence is cited. The monograph is not perfectly evidence-based by its own standards—it is a foundational proposal, not a completed scientific treatise—but it demonstrates the method it advocates.

16.8 On the Peer Review Proposal

The reviewer suggests that the proposal would "effectively exclude" work in aesthetics, phenomenology, political philosophy, and ethics. We disagree. The proposal requires that such work meet minimum standards of formal coherence and empirical compatibility—not that it be reducible to formal models. A phenomenological analysis that is internally consistent, that acknowledges relevant empirical findings, and that specifies conditions of application satisfies the requirements. The bar is not formalization of all content but accountability in method. We do not seek a disciplinary purge; we seek a raising of minimum standards.

17 Limits and Boundaries

Evidence-based philosophy cannot: (i) dissolve metaphysics—it restructures rather than eliminates such questions; (ii) derive values from facts—it constrains normative reasoning through feasibility; (iii) replace philosophical creativity—it constrains evaluation, not generation; (iv) achieve finality—it demands accountability, not certainty. Creativity precedes formalization; formalization disciplines creativity.

18 Broader Research Program

This monograph is part of a broader research program investigating the formal structure of interdisciplinary knowledge systems. Prior work has developed the Reflexive Inference Law establishing information-theoretic bounds on self-referential generalization [61], the concept of Conceptual Responsibility for the ethics of idea transmission [62], and formal frameworks for understanding predictive processing as a physical constraint on adaptive systems [60]. The present monograph addresses a specific gap: the absence of a formal methodology that would allow philosophy to participate as an equal partner in interdisciplinary knowledge synthesis. Without such a methodology, philosophy either remains isolated or is absorbed by the sciences. Evidence-based philosophy provides the missing link: a framework preserving philosophy's distinctive domain while subjecting its methods to standards required for genuine interdisciplinary integration.

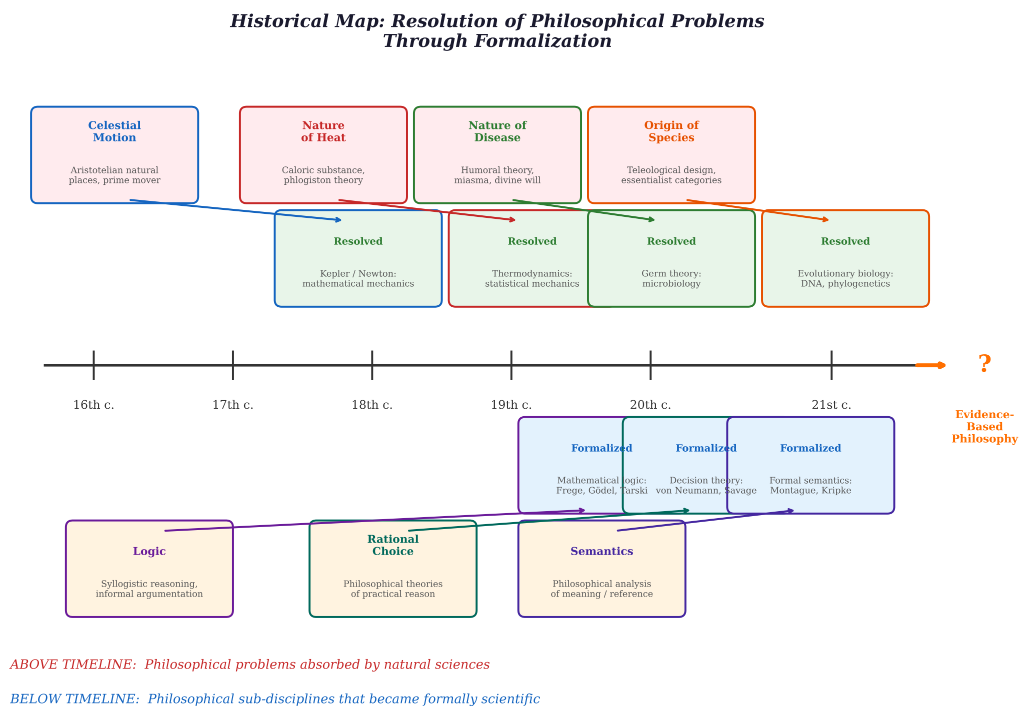

Figure 5: Historical map of philosophical problems resolved through formalization.

19 Conclusion

This monograph has developed a proposal for the structural transformation of philosophy into an evidence-based discipline. The argument rests on three pillars: the demonstrated maturity of formal, empirical, and computational tools; the historical pattern by which formalization resolves philosophical problems; and the emergence of AI as a catalyst removing the cognitive-capacity barrier.

The Reflexive Inference Constraint (Section 11) establishes that evidence-based methods are not merely desirable but informationally necessary for philosophical claims with universal scope: no amount of introspective depth can substitute for the breadth that comparative evidence provides. The Conceptual Responsibility framework (Section 12) shows that formalization serves ethical as well as epistemic functions. The application to predictive processing (Section 9) demonstrates that the method applies to open problems, not only to historically resolved ones.

The proposal is not that philosophy should be abolished or absorbed. It is that philosophy can become a more powerful version of itself—retaining its questions, acquiring more effective methods. The tools are available. The results they can produce are demonstrably superior to those of informal speculation alone. The invitation is to use them.

References

[1] G. Frege, *Begriffsschrift*. Halle: L. Nebert, 1879.

[2] B. Russell and A. N. Whitehead, *Principia Mathematica*, 3 vols. Cambridge: Cambridge University Press, 1910–1913.

[3] R. Carnap, H. Hahn, and O. Neurath, *Wissenschaftliche Weltauffassung: Der Wiener Kreis*. Vienna: Artur Wolf Verlag, 1929.

[4] A. Tarski, "The concept of truth in formalized languages," in *Logic, Semantics, Metamathematics*, pp. 152–278. Oxford: Clarendon Press, 1956 [1933].

[5] R. Montague, "Universal grammar," *Theoria*, vol. 36, pp. 373–398, 1970.

[6] P. W. Bridgman, *The Logic of Modern Physics*. New York: Macmillan, 1927.

[7] K. R. Popper, *The Logic of Scientific Discovery*. London: Hutchinson, 1959.

[8] R. Descartes, *Discourse on the Method*. Leiden, 1637.

[9] F. Bacon, *Novum Organum*. London, 1620.

[10] D. Kahneman and A. Tversky, "Prospect theory," *Econometrica*, vol. 47, no. 2, pp. 263–291, 1979.

[11] L. Wittgenstein, *Tractatus Logico-Philosophicus*. London: Kegan Paul, 1922.

[12] K. Gödel, "Über formal unentscheidbare Sätze," *Monatshefte für Mathematik und Physik*, vol. 38, pp. 173–198, 1931.

[13] A. Church, "An unsolvable problem of elementary number theory," *American Journal of Mathematics*, vol. 58, no. 2, pp. 345–363, 1936.

[14] A. M. Turing, "On computable numbers," *Proc. London Math. Soc.*, vol. 42, pp. 230–265, 1936.

[15] A. J. Ayer, *Language, Truth and Logic*. London: Victor Gollancz, 1936.

[16] I. Lakatos, *The Methodology of Scientific Research Programmes*. Cambridge: Cambridge University Press, 1978.

[17] T. S. Kuhn, *The Structure of Scientific Revolutions*. Chicago: University of Chicago Press, 1962.

[18] P. K. Feyerabend, *Against Method*. London: New Left Books, 1975.

[19] W. V. O. Quine, "Epistemology naturalized," in *Ontological Relativity and Other Essays*, pp. 69–90. New York: Columbia University Press, 1969.

[20] H. Kornblith, *Knowledge and its Place in Nature*. Oxford: Oxford University Press, 2002.

[21] P. Maddy, *Second Philosophy*. Oxford: Oxford University Press, 2007.

[22] J. Knobe, "Intentional action and side effects," *Analysis*, vol. 63, no. 3, pp. 190–194, 2003.

[23] S. Nichols, *Sentimental Rules*. Oxford: Oxford University Press, 2004.

[24] J. M. Weinberg, S. Nichols, and S. Stich, "Normativity and epistemic intuitions," *Philosophical Topics*, vol. 29, pp. 429–460, 2001.

[25] E. Machery, R. Mallon, S. Nichols, and S. Stich, "Semantics, cross-cultural style," *Cognition*, vol. 92, pp. B1–B12, 2004.

[26] S. A. Kripke, *Naming and Necessity*. Cambridge, MA: Harvard University Press, 1980.

[27] H. Leitgeb, *The Stability of Belief*. Oxford: Oxford University Press, 2017.

[28] T. Williamson, *Knowledge and its Limits*. Oxford: Oxford University Press, 2000.

[29] R. Pettigrew, *Accuracy and the Laws of Credence*. Oxford: Oxford University Press, 2016.

[30] K. Easwaran, "Bayesianism I," *Philosophy Compass*, vol. 6, no. 5, pp. 312–320, 2011.

[31] E. Husserl, *Logical Investigations*, 2 vols. London: Routledge, 1900–1901.

[32] M. Heidegger, *Being and Time*. Halle: Max Niemeyer Verlag, 1927.

[33] M. Merleau-Ponty, *Phenomenology of Perception*. Paris: Gallimard, 1945.

[34] H.-G. Gadamer, *Truth and Method*. Tübingen: J. C. B. Mohr, 1960.

[35] J. Habermas, *The Theory of Communicative Action*, 2 vols. Boston: Beacon Press, 1981–1987.

[36] M. Horkheimer and T. W. Adorno, *Dialectic of Enlightenment*. New York: Herder and Herder, 1944.

[37] J. Derrida, *Of Grammatology*. Paris: Les Éditions de Minuit, 1967.

[38] M. Foucault, *The Order of Things*. Paris: Gallimard, 1966.

[39] Dignāga, *Pramāṇa-samuccaya*, c. 5th century CE. See M. Hattori, *Dignāga, On Perception*. Harvard University Press, 1968.

[40] Gautama, *Nyāya Sūtras*, c. 2nd century CE. See S. C. Vidyabhusana, *A History of Indian Logic*. Calcutta University, 1921.

[41] The Mohist Canon, c. 4th century BCE. See A. C. Graham, *Later Mohist Logic, Ethics and Science*. Chinese University Press, 1978.

[42] E. Schwitzgebel, "Do ethicists steal more books?" *Philosophical Psychology*, vol. 22, no. 6, pp. 711–725, 2009.

[43] H. Mercier and D. Sperber, "Why do humans reason?" *Behavioral and Brain Sciences*, vol. 34, no. 2, pp. 57–74, 2011.

[44] L. Floridi, *The Philosophy of Information*. Oxford: Oxford University Press, 2011.

[45] N. Bostrom, *Superintelligence*. Oxford: Oxford University Press, 2014.

[46] D. J. Chalmers, "The singularity: A philosophical analysis," *Journal of Consciousness Studies*, vol. 17, pp. 7–65, 2010.

[47] P. E. Oppenheimer and E. N. Zalta, "A computationally-discovered simplification of the ontological argument," *Australasian Journal of Philosophy*, vol. 89, no. 2, pp. 333–349, 2011.

[48] L. Wittgenstein, *Philosophical Investigations*. Oxford: Blackwell, 1953.

[49] J. von Neumann and O. Morgenstern, *Theory of Games and Economic Behavior*. Princeton University Press, 1944.

[50] L. J. Savage, *The Foundations of Statistics*. New York: Wiley, 1954.

[51] K. J. Arrow, *Social Choice and Individual Values*. New York: Wiley, 1951.

[52] F. de Waal, *Primates and Philosophers*. Princeton University Press, 2006.

[53] M. Tomasello, *Why We Cooperate*. Cambridge, MA: MIT Press, 2009.

[54] R. H. Thaler and C. R. Sunstein, *Nudge*. New Haven: Yale University Press, 2008.

[55] D. Hume, *A Treatise of Human Nature*. London: John Noon, 1739.

[56] I. Kant, *Groundwork of the Metaphysics of Morals*. Riga, 1785.

[57] J. D. Greene, *Moral Tribes*. New York: Penguin Press, 2013.

[58] J. Haidt, *The Righteous Mind*. New York: Vintage Books, 2012.

[59] F. Cushman, "Action, outcome, and value," *Personality and Social Psychology Review*, vol. 17, no. 3, pp. 273–292, 2013.

[60] B. Kriger, "The evolutionary inevitability of predictive processing: A physical constraint argument," Institute of Integrative and Interdisciplinary Research, 2026. doi:10.5281/zenodo.18444910.

[61] B. Kriger, "The Reflexive Inference Law: Bounded generalization from self-observation in inferential systems," Institute of Integrative and Interdisciplinary Research, 2026. doi:10.5281/zenodo.18355847.

[62] B. Kriger, "Conceptual Responsibility: Structural constraints on idea transmission and the ethics of intellectual communication," Information Physics Institute, 2014. doi:10.5281/zenodo.18488818.

[63] D. M. Wolpert and Z. Ghahramani, "Computational principles of movement neuroscience," *Nature Neuroscience*, vol. 3, pp. 1212–1217, 2000.

[64] K. Friston, "The free-energy principle: A unified brain theory?" *Nature Reviews Neuroscience*, vol. 11, no. 2, pp. 127–138, 2010.

[65] F. C. Bartlett, *Remembering: A Study in Experimental and Social Psychology*. Cambridge University Press, 1932.

[66] H. L. Roediger et al., "Bartlett revisited," *Journal of Applied Research in Memory and Cognition*, vol. 3, no. 4, pp. 266–271, 2014.

[67] T. L. Griffiths and M. L. Kalish, "Language evolution by iterated learning with Bayesian agents," *Cognitive Science*, vol. 31, no. 3, pp. 441–480, 2007.

[68] S. Kirby, T. Griffiths, and K. Smith, "Iterated learning and the evolution of language," *Current Opinion in Neurobiology*, vol. 28, pp. 108–114, 2014.

[69] H. Jonas, *The Imperative of Responsibility*. University of Chicago Press, 1984.

[70] T. M. Cover and J. A. Thomas, *Elements of Information Theory*. Wiley, 1991.

[71] E. Schwitzgebel, “The unreliability of naive introspection,” *Philosophical Review*, vol. 117, no. 2, pp. 245–273, 2008.

[72] R. E. Nisbett and T. D. Wilson, “Telling more than we can know,” *Psychological Review*, vol. 84, pp. 231–259, 1977.

[73] J. Henrich, S. J. Heine, and A. Norenzayan, “The weirdest people in the world?” *Behavioral and Brain Sciences*, vol. 33, pp. 61–83, 2010.

[74] D. Parfit, *Reasons and Persons*. Oxford: Clarendon Press, 1984.

[75] C. M. Korsgaard, *The Sources of Normativity*. Cambridge University Press, 1996.

[76] P. F. Strawson, “Freedom and resentment,” *Proceedings of the British Academy*, vol. 48, pp. 1–25, 1962.